Generative Models in Finance

Week 4: Flow Matching and Diffusion Models — Foundations

Cristopher Salvi
Imperial College London

Spring 2026

Overview

1. Generative Modelling as Sampling
2. Flow Models
3. Diffusion Models
4. Constructing the Training Target
5. Training the Generative Model
6. A Guide to the Diffusion Model Literature
7. Building an Image Generator

Reference: P. Holderrieth & E. Erives, *An Introduction to Flow Matching and Diffusion Models*, MIT 6.S184, arXiv:2506.02070, 2025.

Part 1: Generative Modelling as Sampling

- The diversity of valid objects is modelled by a data distribution p_{data} over $\mathbb{R}^d$ (running assumption in these slides)
- Generation = Sampling: generating z means sampling $z \sim p_{\text{data}}$
- Dataset: i.i.d. samples $z_1, \ldots, z_N \sim p_{\text{data}}$
- Conditional generation: sample $z \sim p_{\text{data}}(\cdot \mid y)$ for a conditioning variable y (e.g. text prompt). We focus on unconditional generation first
- Goal: Learn to transform samples from $p_{\text{init}} = \mathcal{N}(0, I_d)$ into samples from p_{data}, via simulation of a suitably constructed *differential equation*

Part 2: Flow Models — ODEs

- A trajectory is a function $X : [0,1] \to \mathbb{R}^d$, $t \mapsto X_t$
- A vector field is a function

$$u : \mathbb{R}^d \times [0,1] \to \mathbb{R}^d, \quad (x,t) \mapsto u_t(x),$$

  specifying a velocity at every point x and time t
- An ordinary differential equation (ODE) with initial condition is:

$$\frac{\mathrm{d}}{\mathrm{d}t} X_t = u_t(X_t), \tag{4.1}$$

$$X_0 = x_0. \tag{4.2}$$

- Equation (4.1) says: the instantaneous rate of change of X_t is prescribed by the vector field u evaluated at the current position X_t

Flows

- The flow of the ODE is the function

$$\psi : \mathbb{R}^d \times [0,1] \to \mathbb{R}^d, \quad (x_0, t) \mapsto \psi_t(x_0)$$

  that maps each initial condition x_0 to the solution $X_t = \psi_t(x_0)$ at time t
- The flow satisfies:

$$\frac{\mathrm{d}}{\mathrm{d}t}\psi_t(x_0) = u_t(\psi_t(x_0)), \tag{4.3}$$

$$\psi_0(x_0) = x_0. \tag{4.4}$$

- Intuition: Vector fields, ODEs, and flows are three descriptions of the same object: *vector fields define ODEs whose solutions are flows*

Existence and Uniqueness

Theorem 4.1 (Flow existence and uniqueness)

If $u : \mathbb{R}^d \times [0,1] \to \mathbb{R}^d$ is continuously differentiable with bounded derivative, then the ODE (4.1)–(4.2) has a unique solution given by a flow ψ_t. Moreover, ψ_t is a diffeomorphism for all t, i.e. ψ_t is continuously differentiable with continuously differentiable inverse ψ_t^{-1}.

- Since we parameterise $u_t(x)$ with neural networks with smooth activations (e.g. GELU, SiLU), the assumptions are approximately satisfied in practice
- Takeaway: Flows exist, are unique, and are diffeomorphisms — they "warp" space in an invertible way

Example 4.2 (Linear vector field)

Let $u_t(x) = -\theta x$ for $\theta > 0$. Then the flow is

$$\psi_t(x_0) = \exp(-\theta t)\, x_0. \tag{4.5}$$

Simulating an ODE: ODE Solvers

- In general, ψ_t cannot be computed in closed form; we approximate it numerically
- Write $t_k = kh$ for $k = 0, \ldots, n$ with step size $h = n^{-1}$, and let $X_{t_0} = x_0$
- Euler method (order 1): one evaluation of u per step

$$X_{t_{k+1}} = X_{t_k} + h\, u_{t_k}(X_{t_k}) \tag{4.6}$$

- The order of a solver measures accuracy: an order-p method has global error $\mathcal{O}(h^p)$
- Higher-order methods use *multiple evaluations* of u per step to achieve better accuracy for the same step size

ODE Solvers: Higher-Order Methods

- Midpoint method (order 2): evaluate u at the midpoint

$$X_{t_k+h/2} = X_{t_k} + \tfrac{h}{2}\, u_{t_k}(X_{t_k}) \quad \blacktriangleright \text{ half-step}$$
$$X_{t_{k+1}} = X_{t_k} + h\, u_{t_k+h/2}(X_{t_k+h/2}) \quad \blacktriangleright \text{ full step}$$

- Heun's method (order 2): predictor-corrector strategy

$$\tilde{X}_{t_{k+1}} = X_{t_k} + h\, u_{t_k}(X_{t_k}) \quad \blacktriangleright \text{ predictor (Euler)}$$
$$X_{t_{k+1}} = X_{t_k} + \frac{h}{2}\big(u_{t_k}(X_{t_k}) + u_{t_{k+1}}(\tilde{X}_{t_{k+1}})\big) \quad \blacktriangleright \text{ corrector}$$

- Both use 2 evaluations of u per step; same order, different geometric properties

ODE Solvers: Runge–Kutta Methods

- The classical 4th-order Runge–Kutta method (RK4):

$$
\begin{aligned}
k_1 &= u_{t_k}(X_{t_k}), \\
k_2 &= u_{t_k + h/2}\big(X_{t_k} + \tfrac{h}{2}\, k_1\big), \\
k_3 &= u_{t_k + h/2}\big(X_{t_k} + \tfrac{h}{2}\, k_2\big), \\
k_4 &= u_{t_{k+1}}\big(X_{t_k} + h\, k_3\big),
\end{aligned}
$$

$$
X_{t_{k+1}} = X_{t_k} + \frac{h}{6}\big(k_1 + 2k_2 + 2k_3 + k_4\big)
$$

- Uses 4 evaluations of u per step; global error $\mathcal{O}(h^4)$
- Euler, midpoint, and Heun are all special cases of the general family of Runge–Kutta methods, characterised by their Butcher tableau

ODE Solvers: Practical Considerations

- Cost–accuracy trade-off: each step of an order-p method requires $s \geq p$ evaluations of u, but fewer steps are needed for the same accuracy
- Adaptive step-size control: modern solvers (e.g. Dormand–Prince / `dopri5`) *adapt* h at each step by comparing two estimates of different orders
- For the rest of these lectures, we use the Euler method for simplicity, but all results generalise to higher-order solvers

Flow Models

- A flow model is defined by:

$$X_0 \sim p_{\text{init}} \qquad \blacktriangleright \text{ random initialisation}$$

$$\frac{\mathrm{d}}{\mathrm{d}t} X_t = u_t^\theta(X_t) \qquad \blacktriangleright \text{ ODE}$$

  where $u_t^\theta : \mathbb{R}^d \times [0,1] \to \mathbb{R}^d$ is a neural network with parameters θ

- Goal: Choose θ so that the endpoint X_1 has distribution p_{data}, i.e.

$$X_1 \sim p_{\text{data}} \quad \Leftrightarrow \quad \psi_1^\theta(X_0) \sim p_{\text{data}}$$

- Important: The neural network parameterises the *vector field*, not the flow itself. The flow is obtained by simulating the ODE

Sampling from a Flow Model

Algorithm 1: Sampling from a Flow Model (Euler method)

Require: Neural network vector field u_t^θ, number of steps n

1. Set $t = 0$, step size $h = 1/n$
2. Draw $X_0 \sim p_{\text{init}}$
3. For $i = 1, \ldots, n$ do:
4. $\quad X_{t+h} = X_t + h\, u_t^\theta(X_t)$
5. $\quad t \leftarrow t + h$
6. Return X_1

Part 3: Diffusion Models — Brownian Motion

Definition 4.3 (Brownian motion)

A Brownian motion (or Wiener process) $W = (W_t)_{0 \le t \le 1}$ is a stochastic process such that $W_0 = 0$, the trajectories $t \mapsto W_t$ are continuous, and:

1. *Normal increments:* $W_t - W_s \sim \mathcal{N}(0, (t-s) I_d)$ *for all* $0 \le s < t$
2. *Independent increments: For any* $0 \le t_0 < t_1 < \cdots < t_n$, *the increments* $W_{t_1} - W_{t_0}, \ldots, W_{t_n} - W_{t_{n-1}}$ *are independent*

- Approximate simulation: set $W_0 = 0$ and update

$$W_{t+h} = W_t + \sqrt{h}\, \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(0, I_d) \tag{4.7}$$

- Brownian paths are continuous but nowhere differentiable (a.s.)
- Brownian motion is the fundamental building block of stochastic calculus, analogous to the Gaussian distribution in probability theory

Advanced BM Simulation: Brownian Bridge and Virtual Brownian Tree

- The sequential scheme (4.7) generates W on a *fixed* grid; adaptive solvers may query W at *arbitrary* times
- Brownian bridge interpolation: Given W_s and W_t, for any $s < r < t$:

$$W_r \mid W_s, W_t \sim \mathcal{N}\Big(\frac{(t-r)W_s + (r-s)W_t}{t-s}, \frac{(r-s)(t-r)}{t-s} I_d\Big)$$

  This allows generating BM values at new times *conditionally* on already-sampled values
- Virtual Brownian Tree[1]: organise time intervals in a *binary tree*; at each query time r, walk down the tree and apply Brownian bridge interpolation
- The entire path is determined by a single PRNG seed: previously generated values need not be stored $\Rightarrow \mathcal{O}(1)$ memory, $\mathcal{O}(\log(1/\mathrm{tol}))$ time per query
- This enables *adaptive step-size* SDE solvers: the solver chooses h on the fly while maintaining a consistent Brownian path across rejected/accepted steps
- Implementation available in Diffrax[2], a JAX library for numerical differential equation solvers

[1] Kidger, *On Neural Differential Equations*, PhD thesis, Oxford, 2021. Extended by Jelinčič, Foster & Kidger, *Single-seed generation of Brownian paths and integrals for adaptive and high order SDE solvers*, arXiv:2405.06464, 2024.

[2] `https://github.com/patrick-kidger/diffrax`

From ODEs to SDEs

- An ODE $\mathrm{d}X_t = u_t(X_t)\,\mathrm{d}t$ defines a *deterministic* flow; given $X_0 = x_0$, the trajectory is uniquely determined
- An Itô stochastic differential equation perturbs this with a martingale driving term:

$$\mathrm{d}X_t = u_t(X_t)\,\mathrm{d}t + \sigma_t\,\mathrm{d}W_t, \tag{4.8}$$

$$X_0 = x_0, \tag{4.9}$$

  where $\sigma_t \geq 0$ is the diffusion coefficient and W is a d-dimensional Brownian motion
- The integral form makes the meaning precise: for all $0 \leq s < t \leq 1$,

$$X_t = X_s + \int_s^t u_r(X_r)\,\mathrm{d}r + \int_s^t \sigma_r\,\mathrm{d}W_r \tag{4.10}$$

  where the second integral is an Itô integral (defined as an L^2-limit of Riemann sums against W)
- Key distinction from ODEs: The solution (X_t) is now a *stochastic process*, not a deterministic path. There is no flow map ψ_t; instead, X_t defines a random variable for each t, adapted to the filtration generated by W

SDE Existence and Uniqueness

Theorem 4.4 (Strong existence and pathwise uniqueness[a])

[a] Mao, X. *Stochastic Differential Equations and Applications*, 2nd ed., Horwood, 2007, Ch. 3.

Suppose that:

1. *(Lipschitz continuity) There exists $L > 0$ such that for all $x, y \in \mathbb{R}^d$ and $t \in [0,1]$:*
$$\|u_t(x) - u_t(y)\| \leq L\|x - y\|$$
2. *(Linear growth) There exists $K > 0$ such that for all $x \in \mathbb{R}^d$ and $t \in [0,1]$:*
$$\|u_t(x)\|^2 + |\sigma_t|^2 \leq K^2(1 + \|x\|^2)$$

Then the SDE (4.8)–(4.9) admits a unique strong solution: a continuous, $(\mathcal{F}_t^W)$-adapted process $(X_t)_{0\leq t\leq 1}$ satisfying (4.10) a.s., with $\mathbb{E}[\sup_{0\leq t\leq 1} \|X_t\|^2] < \infty$.

- The proof proceeds by Picard iteration on the integral equation (4.10), using the Itô isometry and the Borel–Cantelli lemma to upgrade L^2-convergence to a.s. uniform convergence
- Every ODE is a special case with $\sigma_t = 0$. Henceforth, when we speak about SDEs, we consider ODEs as a special case

The Ornstein–Uhlenbeck Process

Example 4.5 (Ornstein–Uhlenbeck process)

Let $u_t(x) = -\theta x$ for $\theta > 0$ and $\sigma_t = \sigma \geq 0$ constant. The SDE

$$\mathrm{d}X_t = -\theta X_t \,\mathrm{d}t + \sigma \,\mathrm{d}W_t \tag{4.11}$$

defines an Ornstein–Uhlenbeck (OU) process.

- The drift is globally Lipschitz with $L = \theta$, so the theorem applies
- Closed-form solution (by variation of constants / integrating factor $e^{\theta t}$):

$$X_t = e^{-\theta t} x_0 + \sigma \int_0^t e^{-\theta(t-s)} \,\mathrm{d}W_s$$

- Since the Itô integral of a deterministic integrand is Gaussian:

$$X_t \sim \mathcal{N}\left(e^{-\theta t} x_0,\ \frac{\sigma^2}{2\theta}(1 - e^{-2\theta t})\, I_d \right)$$

- As $t \to \infty$: $X_t \xrightarrow{d} \mathcal{N}(0, \frac{\sigma^2}{2\theta} I_d)$ (stationary distribution)
- For $\sigma = 0$: reduces to the deterministic flow $\psi_t(x_0) = e^{-\theta t} x_0$ from (4.5)

SDE Solvers: Euler–Maruyama Method

- The Euler–Maruyama method[3] is the SDE analogue of the Euler method for ODEs:

$$X_{t_{k+1}} = X_{t_k} + h\, u_{t_k}(X_{t_k}) + \sigma_{t_k}(W_{t_{k+1}} - W_{t_k}) \tag{4.12}$$

  Since $W_{t_{k+1}} - W_{t_k} \sim \mathcal{N}(0, h\, I_d)$, equivalently: $X_{t_{k+1}} = X_{t_k} + h\, u_{t_k}(X_{t_k}) + \sqrt{h}\, \sigma_{t_k}\, \epsilon_k$, $\epsilon_k \sim \mathcal{N}(0, I_d)$

- When $\sigma_t = 0$, this reduces to the Euler method (4.6)
- Two notions of convergence for SDE solvers:
  - Strong order p: $\mathbb{E}[\|X_T - X_T^h\|] = \mathcal{O}(h^p)$ (pathwise accuracy)
  - Weak order q: $|\mathbb{E}[f(X_T)] - \mathbb{E}[f(X_T^h)]| = \mathcal{O}(h^q)$ for smooth f (distributional accuracy)
- Euler–Maruyama has strong order 1/2 and weak order 1

[3] Maruyama, G. *Continuous Markov processes and stochastic equations*, Rend. Circ. Mat. Palermo, 4:48–90, 1955.

SDE Solvers: Milstein Method

- Euler–Maruyama ignores the interaction between noise and drift. The Milstein method[4] adds a correction term from the Itô–Taylor expansion:

$$X_{t_{k+1}} = X_{t_k} + h\, u_{t_k}(X_{t_k}) + \sigma_{t_k}\, \Delta W_k + \tfrac{1}{2}\, \sigma_{t_k}\, \sigma'_{t_k} \big((\Delta W_k)^2 - h\, I_d \big)$$

  where $\Delta W_k = W_{t_{k+1}} - W_{t_k}$ and $\sigma'_{t_k} = \frac{\mathrm{d}\sigma}{\mathrm{d}t}\big|_{t_k}$

- This achieves strong order 1 (and weak order 1), matching the ODE Euler method's convergence rate
- Caveat (multidimensional case): When $d > 1$ and the diffusion coefficient depends on x, the Milstein correction requires simulating Lévy areas $\int_{t_k}^{t_{k+1}} (W_s^i - W_{t_k}^i)\, \mathrm{d}W_s^j$, which is non-trivial
- For scalar noise ($d = 1$) or additive noise (σ_t independent of x), Milstein is straightforward and recommended over Euler–Maruyama

[4] Milstein, G.N. *Approximate integration of stochastic differential equations*, Theory Probab. Appl., 19(3):557–562, 1974.

SDE Solvers: Higher-Order and Adaptive Methods

- Stochastic Runge–Kutta methods[5]: Achieve strong order 1 (or higher) *without* explicit derivatives of u or σ, by evaluating the coefficients at multiple intermediate points — analogous to deterministic RK methods
- Key difficulty vs. ODEs: Achieving strong order > 1 requires simulating iterated stochastic integrals (Lévy areas), which adds computational cost; most practical SDE solvers therefore operate at strong order ≤ 1
- Adaptive step-size control: Possible via the Virtual Brownian Tree (cf. earlier slide); reject steps where the local error estimate exceeds a tolerance, and retry with smaller h
- For *weak* approximation (sufficient for sampling in generative models), higher weak-order methods exist that avoid Lévy areas entirely

General reference: Kloeden, P.E. & Platen, E. *Numerical Solution of Stochastic Differential Equations*, Springer, 1992.

[5] Rößler, A. *Runge–Kutta methods for the strong approximation of solutions of SDEs*, SIAM J. Numer. Anal., 48(3):922–952, 2010.

The Rough Path Perspective on SDEs

- Core problem: The Itô integral $\int_s^t f(X_r)\,\mathrm{d}W_r$ is defined via L^2-limits. This *probabilistic* construction obscures the *analytic* structure needed for high-order numerics
- Lyons' rough paths theory[6]: Reformulate the SDE as a *pathwise* differential equation driven by the Brownian motion viewed as a rough path
- A rough path over W is not just the path $t \mapsto W_t$, but an enhanced path that also encodes iterated integrals:

$$\mathbf{W}_t = \left(W_t,\ \mathbb{W}_{s,t} := \int_s^t (W_r - W_s) \otimes \circ \mathrm{d}W_r \right)$$

  The second component $\mathbb{W}_{s,t} \in \mathbb{R}^{d \times d}$ is the Lévy area (the antisymmetric part encodes the "winding" of the path). The integral is in the Stratonovich sense
- Key insight: Once $\mathbf{W}$ is fixed, the solution map $\mathbf{W} \mapsto X$ is *continuous* in rough path topology — the Universal Limit Theorem[7]. This yields solutions in the Stratonovich sense; the Itô solution is recovered via the Itô–Stratonovich correction:

$$\int_s^t f(X_r) \circ \mathrm{d}W_r = \int_s^t f(X_r)\,\mathrm{d}W_r + \tfrac{1}{2} \int_s^t f'(X_r)\,\sigma_r\,\mathrm{d}r$$

[6] Lyons, T.J. *Differential equations driven by rough signals*, Rev. Mat. Iberoam., 14(2):215–310, 1998.

[7] Friz, P.K. & Hairer, M. *A Course on Rough Paths*, 2nd ed., Springer, 2020, Ch. 8.

Rough Paths and Numerical Schemes

- The ULT implies: to approximate X, it suffices to approximate the *driving rough path* $\mathbf{W}$ well. Since the ULT operates in the Stratonovich setting, the numerical schemes below approximate the Stratonovich SDE

$$\circ \mathrm{d}X_t = \tilde{u}_t(X_t)\,\mathrm{d}t + \sigma_t \circ \mathrm{d}W_t$$

  where $\tilde{u}_t(x) = u_t(x) - \frac{1}{2}\sigma_t\sigma_t'$ absorbs the Itô–Stratonovich drift correction

- Davie's scheme[8]: Given the Brownian increment $\Delta W_k = W_{t_{k+1}} - W_{t_k}$ and the Lévy area $\mathbb{W}_{t_k,t_{k+1}}$, update via

$$X_{t_{k+1}} = X_{t_k} + \tilde{u}_{t_k}(X_{t_k})\,h + \sigma_{t_k}\,\Delta W_k + D\sigma_{t_k}(X_{t_k})\,\sigma_{t_k}\,\mathbb{W}_{t_k,t_{k+1}}$$

- This achieves strong order 1 and makes explicit *why* Lévy areas are needed: they are part of the rough path data at the required level of regularity
- The Milstein scheme (in Stratonovich form) is recovered as a special case

[8] Davie, A.M. *Differential equations driven by rough paths: an approach via discrete approximation*, Appl. Math. Res. Express, 2008, arXiv:0710.0772.

Log-ODE Method and Higher-Order Rough Path Solvers

- Log-ODE method[9]: On each interval $[t_k, t_{k+1}]$, replace the rough driving signal by a *smooth* path (a geodesic in the free nilpotent group) that matches the same iterated integrals, then solve the resulting ODE exactly or to high order
- This transforms the SDE solve into a sequence of *local ODE solves*, for which classical Runge–Kutta methods apply
- Convergence: Approximating the rough path signature to level N yields a method of strong order $N/2 - \varepsilon$. Level $N = 2$ (increments + Lévy area) gives strong order 1; level $N = 3$ adds further iterated integrals for strong order $3/2$
- Dimension-free error bounds[10]: Rough path estimates give error bounds that depend on the *regularity* of the driving signal, not the ambient dimension d — a crucial advantage for high-dimensional problems in generative modelling

General reference: Friz, P.K. & Victoir, N.B. *Multidimensional Stochastic Processes as Rough Paths*, Cambridge, 2010.

[9] Foster, J., Lyons, T. & Oberhauser, H. *An optimal polynomial approximation of Brownian motion*, SIAM J. Numer. Anal., 58(3):1393–1421, 2020, arXiv:1904.06998.

[10] Boutaib, Y., Gyurkó, L.G., Lyons, T. & Yang, D. *Dimension-free Euler estimates of rough differential equations*, Rev. Roumaine Math. Pures Appl., 59(1):25–53, 2014, arXiv:1307.4708.

Diffusion Models

- A diffusion model is defined by:

$$\mathrm{d}X_t = u_t^\theta(X_t)\,\mathrm{d}t + \sigma_t\,\mathrm{d}W_t \qquad \blacktriangleright \text{ SDE}$$

$$X_0 \sim p_{\mathrm{init}} \qquad \blacktriangleright \text{ random initialisation}$$

- It consists of:
  - A neural network $u^\theta : \mathbb{R}^d \times [0,1] \to \mathbb{R}^d$, $(x,t) \mapsto u_t^\theta(x)$ with parameters θ
  - A fixed diffusion coefficient $\sigma_t : [0,1] \to [0,\infty)$
- Goal: $X_1 \sim p_{\mathrm{data}}$
- A diffusion model with $\sigma_t = 0$ is a flow model

Sampling from a Diffusion Model

Algorithm 2: Sampling from a Diffusion Model (Euler–Maruyama)

Require: Neural network u_t^θ, number of steps n, diffusion coefficient σ_t

1. Set $t = 0$, step size $h = 1/n$
2. Draw $X_0 \sim p_{\text{init}}$
3. For $i = 1, \ldots, n$ do:
4. Draw $\epsilon \sim \mathcal{N}(0, I_d)$
5. $X_{t+h} = X_t + h\, u_t^\theta(X_t) + \sigma_t \sqrt{h}\, \epsilon$
6. $t \leftarrow t + h$
7. Return X_1

Part 4: Constructing the Training Target

- If we randomly initialise θ, simulating the ODE/SDE produces nonsense
- We need to *train* the neural network by minimising a loss:

$$\mathcal{L}(\theta) = \|u_t^\theta(x) - u_t^{\text{target}}(x)\|^2 \tag{4.13}$$

- The training target $u_t^{\text{target}} : \mathbb{R}^d \times [0,1] \to \mathbb{R}^d$ is a vector field whose corresponding ODE/SDE converts p_{init} into p_{data}
- This part: Derive a formula for u_t^{target}
- Next part: Design a tractable training algorithm to approximate u_t^{target}

Conditional Probability Paths

Definition 4.6 (Conditional probability path)

A conditional (interpolating) probability path is a family of distributions $p_t(\cdot \mid z)$ over $\mathbb{R}^d$, parameterised by $t \in [0,1]$ and $z \in \mathbb{R}^d$, such that

$$p_0(\cdot \mid z) = p_{\text{init}}, \qquad p_1(\cdot \mid z) = \delta_z, \qquad \text{for all } z \in \mathbb{R}^d. \tag{4.14}$$

- Here δ_z is the Dirac delta at z (the distribution that always returns z)
- Intuitively, $p_t(\cdot \mid z)$ gradually converts a *single* data point z from noise p_{init} at $t = 0$ to a point mass at z at $t = 1$
- Think of this as a trajectory in the space of distributions, indexed by t

Marginal Probability Paths

Definition 4.7 (Marginal probability path)

Given a conditional probability path $\{p_t(\cdot \mid z)\}_{t\in[0,1],z\in\mathbb{R}^d}$ *and a data distribution* p_{data}*, the marginal probability path* $\{p_t\}_{t\in[0,1]}$ *is the mixture*

$$p_t(x) = \int_{\mathbb{R}^d} p_t(x \mid z)\, p_{\text{data}}(z)\, \mathrm{d}z = \mathbb{E}_{z\sim p_{\text{data}}}\big[p_t(x \mid z)\big]. \tag{4.15}$$

- Equivalent sampling procedure:

$$z \sim p_{\text{data}}, \quad x \sim p_t(\cdot \mid z) \quad \Rightarrow \quad x \sim p_t \tag{4.16}$$

- Boundary conditions (inherited from the conditional path):

$$p_0(x) = \int p_0(x \mid z)\, p_{\text{data}}(z)\, \mathrm{d}z = \int p_{\text{init}}(x)\, p_{\text{data}}(z)\, \mathrm{d}z = p_{\text{init}}(x), \tag{4.17}$$

$$p_1(x) = \int \delta_z(x)\, p_{\text{data}}(z)\, \mathrm{d}z = p_{\text{data}}(x). \tag{4.18}$$

- Computational note: We can *sample* from p_t via (4.16), but we cannot *evaluate* $p_t(x)$ because the integral (4.15) is intractable (it marginalises over the entire data distribution)

Gaussian Conditional Probability Paths

Example 4.8 (Gaussian conditional probability path)

Let α_t, β_t be noise schedulers: continuously differentiable, monotonic functions with $\alpha_0 = \beta_1 = 0$ and $\alpha_1 = \beta_0 = 1$. Define

$$p_t(\cdot \mid z) = \mathcal{N}(\alpha_t z, \beta_t^2 I_d). \tag{4.19}$$

- Check: $p_0(\cdot \mid z) = \mathcal{N}(0, \beta_0^2 I_d) = \mathcal{N}(0, I_d) = p_{\text{init}}$ ✓
- Check: $p_1(\cdot \mid z) = \mathcal{N}(\alpha_1 z, 0) = \delta_z$ ✓
- Sampling from the marginal path:

$$z \sim p_{\text{data}}, \quad \epsilon \sim \mathcal{N}(0, I_d) \quad \Rightarrow \quad x = \alpha_t z + \beta_t \epsilon \sim p_t \tag{4.20}$$

- Intuitively: more noise for smaller t (closer to $t = 0$, pure noise), less noise for larger t (closer to $t = 1$, pure data)

Interlude: The Wasserstein-2 Distance

- Optimal transport[11] asks: what is the cheapest way to move mass from a distribution μ to a distribution ν, when transporting a unit of mass from x to y costs $\|x-y\|^2$?
- The Wasserstein-2 distance solves this by optimising over all couplings (joint distributions with the correct marginals):

$$W_2^2(\mu,\nu) = \inf_{\gamma\in\Pi(\mu,\nu)} \int_{\mathbb{R}^d\times\mathbb{R}^d} \|x-y\|^2 \,\mathrm{d}\gamma(x,y),$$

  where $\Pi(\mu,\nu) = \{\gamma \in \mathcal{P}(\mathbb{R}^d\times\mathbb{R}^d) : \gamma(\cdot,\mathbb{R}^d) = \mu,\ \gamma(\mathbb{R}^d,\cdot) = \nu\}$
- Geodesic: $(\mathcal{P}_2(\mathbb{R}^d), W_2)$ is a metric space[12]. A constant-speed geodesic is a curve $(\mu_t)_{t\in[0,1]}$ satisfying

$$W_2(\mu_s,\,\mu_t) = |t-s|\,W_2(\mu_0,\,\mu_1), \qquad \forall\, s,t\in[0,1].$$

- Displacement interpolation[13]: given an optimal transport map T, the path $\mu_t = \big((1-t)\,\mathrm{id} + t\,T\big)_{\#}\mu$ is the unique constant-speed geodesic from $\mu_0=\mu$ to $\mu_1 = T_{\#}\mu = \nu$

[11] Villani, C. (2009). *Optimal Transport: Old and New*. Springer. Grundlehren der mathematischen Wissenschaften, Vol. 338.

[12] Ambrosio, L., Gigli, N., Savaré, G. (2008). *Gradient Flows in Metric Spaces and in the Space of Probability Measures*. 2nd ed., Birkhäuser. Lectures in Mathematics ETH Zürich.

[13] McCann, R.J. (1997). A convexity principle for interacting gases. *Advances in Mathematics*, 128(1), 153–179.

The Conditional Optimal Transport Path

Set $\alpha_t = t$, $\beta_t = 1 - t$[14]:

$$p_t(\cdot \mid z) = \mathcal{N}\big(t\,z,\ (1-t)^2 I_d\big), \qquad X_t = t\,z + (1-t)\,\epsilon, \quad \epsilon \sim \mathcal{N}(0, I_d). \quad (4.21)$$

- The unique OT map from $\mathcal{N}(0, I_d)$ to δ_z is $T(x) = z$ (all mass goes to z)
- The displacement interpolation $T_t(x) = (1-t)\,x + t\,z$ recovers exactly (4.21)
- Hence the CondOT path is the *W_2-geodesic* from $\mathcal{N}(0, I_d)$ to δ_z
- Key consequence: Each sample follows a *straight line* with constant velocity:

$$\dot{X}_t = z - \epsilon$$

Straighter trajectories $\Rightarrow$ fewer ODE discretisation steps at inference

[14] Lipman, Y., Chen, R.T.Q., Ben-Hamu, H., Nickel, M. (2023). Flow Matching for Generative Modeling. *ICLR 2023*. arXiv:2210.02747.

Conditional Vector Fields

- We now construct a vector field whose ODE trajectories follow the conditional probability path
- A conditional vector field $u_t^{\text{target}}(\cdot \mid z)$ is defined so that its ODE yields $X_t \sim p_t(\cdot \mid z)$:

$$X_0 \sim p_{\text{init}}, \quad \frac{\mathrm{d}}{\mathrm{d}t} X_t = u_t^{\text{target}}(X_t \mid z) \quad \Rightarrow \quad X_t \sim p_t(\cdot \mid z). \tag{4.22}$$

Conditional Gaussian Vector Field

Example 4.9 (Target ODE for Gaussian probability paths)

For $p_t(\cdot \mid z) = \mathcal{N}(\alpha_t z, \beta_t^2 I_d)$, the conditional Gaussian vector field is

$$u_t^{\text{target}}(x \mid z) = \left(\dot{\alpha}_t - \frac{\dot{\beta}_t}{\beta_t} \alpha_t \right) z + \frac{\dot{\beta}_t}{\beta_t} x. \tag{4.23}$$

Proof. Define the map $\psi_t^{\text{target}}(x \mid z) = \alpha_t z + \beta_t x$. This is a valid flow:

- $\psi_0^{\text{target}}(x \mid z) = \alpha_0 z + \beta_0 x = 0 + 1 \cdot x = x \quad (\psi_0 = \text{id})$ ✓
- For $t \in [0, 1)$, $\beta_t > 0$, so $\psi_t^{\text{target}}(\cdot \mid z)$ is a diffeomorphism with inverse $(\psi_t^{\text{target}})^{-1}(y \mid z) = (y - \alpha_t z)/\beta_t$ ✓

If $X_0 \sim \mathcal{N}(0, I_d) = p_{\text{init}}$, then by the affine property of Gaussians:

$$X_t = \psi_t^{\text{target}}(X_0 \mid z) = \alpha_t z + \beta_t X_0 \sim \mathcal{N}(\alpha_t z, \beta_t^2 I_d) = p_t(\cdot \mid z). \tag{4.24}$$

Conditional Gaussian Vector Field (Proof cont.)

It remains to extract $u_t^{\text{target}}(x \mid z)$ from ψ_t^{target}. By definition of a flow (4.3):

$$\frac{\mathrm{d}}{\mathrm{d}t}\psi_t^{\text{target}}(x \mid z) = u_t^{\text{target}}(\psi_t^{\text{target}}(x \mid z) \mid z) \quad \text{for all } x, z \in \mathbb{R}^d.$$

Differentiating $\psi_t^{\text{target}}(x \mid z) = \alpha_t z + \beta_t x$ with respect to t:

$$\dot{\alpha}_t z + \dot{\beta}_t x = u_t^{\text{target}}(\alpha_t z + \beta_t x \mid z).$$

Since $\beta_t > 0$ for $t \in [0, 1)$, the map $x \mapsto \alpha_t z + \beta_t x$ is surjective onto $\mathbb{R}^d$. Substituting $x = (\psi_t^{\text{target}})^{-1}(y \mid z) = (y - \alpha_t z)/\beta_t$:

$$\dot{\alpha}_t z + \dot{\beta}_t \cdot \frac{y - \alpha_t z}{\beta_t} = u_t^{\text{target}}(y \mid z) \quad \text{for all } y \in \mathbb{R}^d.$$

Rearranging (collecting the z and y terms):

$$u_t^{\text{target}}(y \mid z) = \left(\dot{\alpha}_t - \frac{\dot{\beta}_t}{\beta_t}\alpha_t\right) z + \frac{\dot{\beta}_t}{\beta_t} y. \qquad \square$$

The Marginalization Trick

Theorem 4.10 (Marginalization trick)

Let $u_t^{\text{target}}(\cdot \mid z)$ be a conditional vector field satisfying (4.22). Then the marginal vector field

$$u_t^{\text{target}}(x) = \int u_t^{\text{target}}(x \mid z) \frac{p_t(x \mid z)\, p_{\text{data}}(z)}{p_t(x)} \, \mathrm{d}z \tag{4.25}$$

yields ODE trajectories that follow the marginal probability path:

$$X_0 \sim p_{\text{init}}, \quad \frac{\mathrm{d}}{\mathrm{d}t} X_t = u_t^{\text{target}}(X_t) \quad \Rightarrow \quad X_t \sim p_t \quad (0 \le t \le 1). \tag{4.26}$$

In particular, $X_1 \sim p_{\text{data}}$.

- Key insight: We can construct the (intractable) marginal vector field from the (tractable) conditional vector field via (4.25)
- The proof uses the continuity equation

The Continuity Equation

Theorem 4.11 (Continuity equation)

Let u_t^{target} *be a vector field with* $X_0 \sim p_{\text{init}}$. *Then* $X_t \sim p_t$ *for all* $0 \leq t \leq 1$ *if and only if*

$$\partial_t p_t(x) = -\text{div}(p_t\, u_t^{\text{target}})(x) \quad \text{for all } x \in \mathbb{R}^d,\ 0 \leq t \leq 1, \tag{4.27}$$

where $\text{div}(v_t)(x) = \sum_{i=1}^d \frac{\partial}{\partial x_i} v_t^i(x)$ *is the divergence operator.*

- Interpretation: The left side $\partial_t p_t(x)$ measures how probability density changes at x. The right side $-\text{div}(p_t u_t)$ measures the net inflow of probability mass via the vector field. Since mass is conserved, these must be equal
- This is a special case ($\sigma_t = 0$) of the Fokker–Planck equation (Theorem 4.14), whose proof is given later. We use this result to prove the Marginalization Trick

Proof of the Marginalization Trick

We show that $u_t^{\text{target}}(x)$ defined in (4.25) satisfies the continuity equation (4.27).

$$
\begin{aligned}
\partial_t p_t(x) &\overset{(i)}{=} \partial_t \int p_t(x \mid z)\, p_{\text{data}}(z)\, \mathrm{d}z \\
&= \int \partial_t p_t(x \mid z)\, p_{\text{data}}(z)\, \mathrm{d}z \\
&\overset{(ii)}{=} \int -\mathrm{div}\big(p_t(\cdot \mid z)\, u_t^{\text{target}}(\cdot \mid z)\big)(x)\, p_{\text{data}}(z)\, \mathrm{d}z \\
&\overset{(iii)}{=} -\mathrm{div}\left(\int p_t(x \mid z)\, u_t^{\text{target}}(x \mid z)\, p_{\text{data}}(z)\, \mathrm{d}z\right)
\end{aligned}
$$

where:

- (i) definition of $p_t(x)$ as $\int p_t(x \mid z)\, p_{\text{data}}(z)\, \mathrm{d}z$ (equation (4.15));
- (ii) the conditional vector field satisfies the continuity equation: $\partial_t p_t(\cdot \mid z) = -\mathrm{div}(p_t(\cdot \mid z)\, u_t^{\text{target}}(\cdot \mid z))$ (this holds by hypothesis (4.22) and Theorem 4.11 applied to each fixed z);
- (iii) interchange of div_x and $\int \mathrm{d}z$ (justified by Fubini–Tonelli, using the regularity of u_t^{target} and p_t). Passing ∂_t under the integral from line 1 to line 2 is similarly justified by dominated convergence.

Proof of the Marginalization Trick (cont.)

Continuing:

$$\overset{(iv)}{=} -\operatorname{div}\left(p_t(x)\int u_t^{\text{target}}(x \mid z)\,\frac{p_t(x \mid z)\,p_{\text{data}}(z)}{p_t(x)}\,\mathrm{d}z\right)$$
$$\overset{(v)}{=} -\operatorname{div}\big(p_t\,u_t^{\text{target}}\big)(x),$$

where:

- (iv) factor out $p_t(x) > 0$ (the marginal density is strictly positive on its support, since it is a mixture of Gaussians);
- (v) recognise the definition of $u_t^{\text{target}}(x)$ in (4.25).

The first and last expressions give $\partial_t p_t(x) = -\operatorname{div}(p_t\,u_t^{\text{target}})(x)$, which is the continuity equation (4.27) for the marginal vector field. By Theorem 4.11, this implies $X_t \sim p_t$ for $0 \le t \le 1$. □

The Score Function

Definition 4.12 (Score function)

The marginal score function of p_t is

$$\nabla \log p_t(x) = \frac{\nabla p_t(x)}{p_t(x)}.$$

The conditional score function of $p_t(\cdot \mid z)$ is $\nabla \log p_t(x \mid z)$.

- The marginal score can be expressed via conditional scores:

$$\nabla \log p_t(x) = \int \nabla \log p_t(x \mid z) \, \frac{p_t(x \mid z) \, p_{\text{data}}(z)}{p_t(x)} \, \mathrm{d}z \tag{4.28}$$

- The conditional score is typically known analytically

Score Function for Gaussian Paths

Example 4.13 (Score for Gaussian conditional probability paths)

For $p_t(x \mid z) = \mathcal{N}(x; \alpha_t z, \beta_t^2 I_d)$, the conditional score is

$$\nabla \log p_t(x \mid z) = -\frac{x - \alpha_t z}{\beta_t^2}. \tag{4.29}$$

- Proof. The Gaussian density is $\mathcal{N}(x; \mu, \sigma^2 I_d) = (2\pi\sigma^2)^{-d/2} \exp\left(-\frac{\|x-\mu\|^2}{2\sigma^2}\right)$. Taking $\nabla_x \log$ gives $-\frac{x-\mu}{\sigma^2}$. With $\mu = \alpha_t z$ and $\sigma^2 = \beta_t^2$, the result follows
- Key property: The conditional score is *linear* in x. This is a special feature of Gaussian distributions

From ODEs to SDEs: Extending the Continuity Equation

- So far, we have derived the training target for flow models (ODEs):
  - Conditional vector field $u_t^{\text{target}}(x \mid z)$ generates the conditional path $p_t(\cdot \mid z)$
  - Marginalization trick (Theorem 4.10): the marginal vector field generates p_t
  - The proof relied on the continuity equation (4.27)
- Goal: Extend this to diffusion models (SDEs) with $\sigma_t > 0$
- Problem: The continuity equation only governs ODEs. For SDEs, we need to account for the additional spreading of probability mass caused by Brownian noise
- Solution: The Fokker–Planck equation — the SDE analogue of the continuity equation. It adds a second-order (diffusion) term $\frac{\sigma_t^2}{2}\Delta p_t$ to capture the effect of noise
- We will use the score function $\nabla \log p_t$ (just defined) to connect the Fokker–Planck equation back to a concrete SDE formula for diffusion models

The Fokker–Planck Equation

Theorem 4.14 (Fokker–Planck equation)

Let $u_t : \mathbb{R}^d \to \mathbb{R}^d$ be Lipschitz continuous and $\sigma_t \geq 0$ continuous. Consider the SDE

$$X_0 \sim p_{\text{init}}, \quad \mathrm{d}X_t = u_t(X_t)\,\mathrm{d}t + \sigma_t\,\mathrm{d}W_t.$$

Suppose the law of X_t admits a smooth density p_t for each $t \in [0,1]$. Then p_t satisfies the Fokker–Planck equation:

$$\partial_t p_t(x) = -\mathrm{div}(p_t u_t)(x) + \frac{\sigma_t^2}{2}\Delta p_t(x), \tag{4.30}$$

where $\Delta p_t(x) = \sum_{i=1}^d \frac{\partial^2}{\partial x_i^2} p_t(x)$ is the Laplacian.

Conversely, if a smooth probability density p_t satisfies (4.30) with $p_0 = p_{\text{init}}$, then $X_t \sim p_t$ for all t.

- For $\sigma_t = 0$: reduces to the continuity equation (4.27)
- The term $\frac{\sigma_t^2}{2}\Delta p_t$ captures the spreading of probability mass due to Brownian noise (cf. the heat equation $\partial_t p = \frac{\sigma^2}{2}\Delta p$, which is the special case $u_t = 0$)

Proof of the Fokker–Planck Equation (Necessity)

Goal: Show that the density p_t of X_t satisfies (4.30). The strategy is to test against an arbitrary $f \in C_c^\infty(\mathbb{R}^d)$ (smooth, compactly supported test function) and derive the PDE in the *weak* (distributional) sense.

Step 1 (Itô's formula). Since f is C^2 and X_t satisfies the SDE $\mathrm{d}X_t = u_t(X_t)\,\mathrm{d}t + \sigma_t\,\mathrm{d}W_t$, Itô's formula gives:

$$\mathrm{d}f(X_t) = \nabla f(X_t)^T \mathrm{d}X_t + \frac{1}{2}\sum_{i,j=1}^{d} \frac{\partial^2 f}{\partial x_i \partial x_j}(X_t)\,[\mathrm{d}X_t]_i\,[\mathrm{d}X_t]_j.$$

The quadratic variation of $\mathrm{d}X_t = u_t\,\mathrm{d}t + \sigma_t\,\mathrm{d}W_t$ is $[\mathrm{d}X_t^i][\mathrm{d}X_t^j] = \sigma_t^2\,\delta_{ij}\,\mathrm{d}t$ (since σ_t is scalar and $\mathrm{d}t\,\mathrm{d}t = \mathrm{d}t\,\mathrm{d}W = 0$, $\mathrm{d}W^i\mathrm{d}W^j = \delta_{ij}\mathrm{d}t$). Therefore:

$$\mathrm{d}f(X_t) = \underbrace{\nabla f(X_t)^T u_t(X_t)}_{\text{drift}}\,\mathrm{d}t + \underbrace{\sigma_t \nabla f(X_t)^T \mathrm{d}W_t}_{\text{martingale}} + \underbrace{\frac{\sigma_t^2}{2}\Delta f(X_t)}_{\text{Itô correction}}\,\mathrm{d}t.$$

Proof of the Fokker–Planck Equation (Necessity, cont.)

Step 2a (Integrate Step 1). Integrate the Itô formula from t to $t+h$:

$$f(X_{t+h}) - f(X_t) = \int_t^{t+h} \nabla f(X_s)^T u_s(X_s)\,\mathrm{d}s + \int_t^{t+h} \sigma_s \nabla f(X_s)^T \mathrm{d}W_s + \int_t^{t+h} \frac{\sigma_s^2}{2} \Delta f(X_s)\,\mathrm{d}s.$$

Step 2b (Take expectation). The three terms are handled as follows:

- *Terms 1 & 3* (Lebesgue integrals): swap $\mathbb{E}$ and $\int \mathrm{d}s$ by Fubini–Tonelli, justified by the boundedness of ∇f, Δf (since $f \in C_c^\infty$) and the moment bound $\mathbb{E}[\sup_s \|X_s\|^2] < \infty$
- *Term 2* (Itô integral): since $f \in C_c^\infty$, the integrand $\sigma_s \nabla f(X_s)$ is bounded, so the Itô integral is a true martingale. Hence $\mathbb{E}\big[\int_t^{t+h} \sigma_s \nabla f(X_s)^T \mathrm{d}W_s\big] = 0$

Proof of the Fokker–Planck Equation (Necessity, cont.)

Step 2b (cont.) After taking expectations, we are left with:

$$\mathbb{E}[f(X_{t+h})] - \mathbb{E}[f(X_t)] = \int_t^{t+h} \mathbb{E}\left[\nabla f(X_s)^T u_s(X_s) + \frac{\sigma_s^2}{2}\Delta f(X_s)\right] \mathrm{d}s.$$

Step 2c (Differentiate). Divide both sides by $h > 0$ and let $h \to 0$. By the fundamental theorem of calculus:

$$\frac{\mathrm{d}}{\mathrm{d}t}\mathbb{E}[f(X_t)] = \mathbb{E}\left[\nabla f(X_t)^T u_t(X_t) + \frac{\sigma_t^2}{2}\Delta f(X_t)\right]. \tag{4.31}$$

Proof of the Fokker–Planck Equation (Necessity, cont.)

Step 3 (Rewrite using the density p_t). Since X_t has density p_t, every expectation can be written as an integral against p_t. Equation (4.31) becomes:

$$\frac{\mathrm{d}}{\mathrm{d}t}\int_{\mathbb{R}^d} f(x)\, p_t(x)\,\mathrm{d}x = \int_{\mathbb{R}^d}\left[\nabla f(x)^T u_t(x) + \frac{\sigma_t^2}{2}\Delta f(x)\right] p_t(x)\,\mathrm{d}x.$$

We pass $\frac{\mathrm{d}}{\mathrm{d}t}$ under the integral on the LHS. This is justified by dominated convergence: f has compact support and hence $|f(x)\,\partial_t p_t(x)| \le \|f\|_\infty \sup_{s\in[0,1]} |\partial_t p_s(x)|$, which is integrable by our smoothness assumption on p_t. We obtain:

$$\int_{\mathbb{R}^d} f(x)\,\partial_t p_t(x)\,\mathrm{d}x = \int_{\mathbb{R}^d} \nabla f(x)^T\big(p_t(x)\,u_t(x)\big)\,\mathrm{d}x + \int_{\mathbb{R}^d} \frac{\sigma_t^2}{2}\,\Delta f(x)\,p_t(x)\,\mathrm{d}x. \tag{4.32}$$

Proof of the Fokker–Planck Equation (Necessity, cont.)

Step 4 (Integration by parts — transfer derivatives from f to p_t).

Recall: the multivariable integration-by-parts formula (a consequence of the divergence theorem) states: for smooth $\varphi : \mathbb{R}^d \to \mathbb{R}$ and $V : \mathbb{R}^d \to \mathbb{R}^d$ with φ compactly supported,

$$\int_{\mathbb{R}^d} \nabla\varphi \cdot V \, \mathrm{d}x = -\int_{\mathbb{R}^d} \varphi \, \mathrm{div}(V) \, \mathrm{d}x. \tag{4.33}$$

No boundary term appears because φ vanishes outside a compact set.

First term of (4.32). Apply (4.33) with $\varphi = f$ and $V = p_t\, u_t$:

$$\int_{\mathbb{R}^d} \nabla f(x)^T (p_t u_t)(x) \, \mathrm{d}x = -\int_{\mathbb{R}^d} f(x) \, \mathrm{div}(p_t u_t)(x) \, \mathrm{d}x. \tag{4.34}$$

Proof of the Fokker–Planck Equation (Necessity, cont.)

Second term of (4.32). We need to move both derivatives from Δf onto p_t. Since $\Delta f = \sum_{i=1}^{d} \partial_{x_i}^2 f$, it suffices to handle each coordinate and sum. We integrate by parts *twice*, one derivative at a time.

1st integration by parts. Apply (4.33) with $\varphi = f$ and $V = p_t \nabla f$ coordinate-wise. Since $\Delta f = \operatorname{div}(\nabla f)$:

$$\int_{\mathbb{R}^d} (\Delta f)\, p_t \, \mathrm{d}x = \int_{\mathbb{R}^d} \operatorname{div}(\nabla f)\, p_t \, \mathrm{d}x = -\int_{\mathbb{R}^d} \nabla f \cdot \nabla p_t \, \mathrm{d}x.$$

(Here (4.33) is applied with $\varphi = p_t$ and $V = \nabla f$; the boundary term vanishes since ∇f is compactly supported.)

2nd integration by parts. Apply (4.33) again, now with $\varphi = f$ and $V = \nabla p_t$:

$$-\int_{\mathbb{R}^d} \nabla f \cdot \nabla p_t \, \mathrm{d}x = +\int_{\mathbb{R}^d} f \operatorname{div}(\nabla p_t) \, \mathrm{d}x = +\int_{\mathbb{R}^d} f \, \Delta p_t \, \mathrm{d}x.$$

Combining: $\displaystyle\int_{\mathbb{R}^d} (\Delta f)\, p_t \, \mathrm{d}x = \int_{\mathbb{R}^d} f \, \Delta p_t \, \mathrm{d}x.$

Proof of the Fokker–Planck Equation (Necessity, cont.)

Step 4 (Conclusion). Substituting (4.34) and the result for the second term into (4.32):

$$\underbrace{\int_{\mathbb{R}^d} f\,\partial_t p_t\,\mathrm{d}x}_{\text{LHS of (4.32)}} = \underbrace{-\int_{\mathbb{R}^d} f\,\mathrm{div}(p_t u_t)\,\mathrm{d}x}_{\text{1st term after IBP}} + \underbrace{\frac{\sigma_t^2}{2}\int_{\mathbb{R}^d} f\,\Delta p_t\,\mathrm{d}x}_{\text{2nd term after IBP}}.$$

Moving everything to one side:

$$\int_{\mathbb{R}^d} f(x)\left[\partial_t p_t(x) + \mathrm{div}(p_t u_t)(x) - \frac{\sigma_t^2}{2}\Delta p_t(x)\right]\mathrm{d}x = 0.$$

This holds for *every* test function $f \in C_c^\infty(\mathbb{R}^d)$.

Proof of the Fokker–Planck Equation (Necessity, cont.)

Step 5 (Conclude). We have shown that for every test function $f \in C_c^\infty(\mathbb{R}^d)$:

$$\int_{\mathbb{R}^d} f(x) \left[\partial_t p_t(x) + \operatorname{div}(p_t u_t)(x) - \frac{\sigma_t^2}{2} \Delta p_t(x) \right] \mathrm{d}x = 0.$$

By the fundamental lemma of the calculus of variations (du Bois-Reymond): if a continuous function integrates to zero against every $f \in C_c^\infty$, it must be identically zero. Therefore:

$$\partial_t p_t(x) = -\operatorname{div}(p_t u_t)(x) + \frac{\sigma_t^2}{2} \Delta p_t(x) \quad \text{for all } x \in \mathbb{R}^d,\ t \in [0,1]. \qquad \square$$

Proof of the Fokker–Planck Equation (Sufficiency)

Claim: If a smooth probability density p_t satisfies (4.30) with $p_0 = p_{\text{init}}$, then $X_t \sim p_t$ for all $t \in [0, 1]$.

Proof. Let q_t denote the density of X_t. By the necessity direction, q_t satisfies FP with $q_0 = p_{\text{init}}$. By hypothesis, p_t satisfies the same PDE with the same initial condition. Since FP is a linear parabolic PDE (with bounded coefficients), it has a unique solution[15,16]. Therefore $p_t = q_t$ for all t. □

[15] Friedman, A. (1964). *Partial Differential Equations of Parabolic Type*. Prentice-Hall, Ch. 2.
[16] Evans, L.C. (2010). *Partial Differential Equations*. 2nd ed., AMS, Ch. 7, §2.3.

From Flow Models to Diffusion Models

Where we stand. The Marginalization Trick (Theorem 4.10) gave us an ODE — a *flow model* — whose trajectories follow the marginal path p_t from p_{init} to p_{data}:

$$\mathrm{d}X_t = u_t^{\text{target}}(X_t)\,\mathrm{d}t.$$

The question. Suppose we want to build a *diffusion model* instead: we add Brownian noise $\sigma_t\,\mathrm{d}W_t$ to the dynamics. The noise perturbs the trajectories, so X_t would no longer follow p_t. *How must we modify the drift to compensate, so that $X_t \sim p_t$ is preserved?*

The SDE Extension Trick

Theorem 4.15 (SDE extension trick)

Let $u_t^{\text{target}}(x)$ be the marginal vector field from Theorem 4.10, whose ODE generates the marginal path p_t. Then for any diffusion coefficient $\sigma_t \geq 0$, the SDE

$$X_0 \sim p_{\text{init}}, \quad \mathrm{d}X_t = \left[u_t^{\text{target}}(X_t) + \frac{\sigma_t^2}{2} \nabla \log p_t(X_t) \right] \mathrm{d}t + \sigma_t \, \mathrm{d}W_t \tag{4.35}$$

satisfies $X_t \sim p_t$ for $0 \leq t \leq 1$. In particular, $X_1 \sim p_{\text{data}}$.

- For $\sigma_t = 0$: recovers the ODE $\mathrm{d}X_t = u_t^{\text{target}}(X_t) \, \mathrm{d}t$
- Key insight: Adding noise ($\sigma_t > 0$) requires a score-based correction $\frac{\sigma_t^2}{2} \nabla \log p_t$ to the drift in order to preserve the same marginal path p_t

Proof of the SDE Extension Trick

What we need to show. The SDE (4.35) has drift $\tilde{u}_t(x) = u_t^{\text{target}}(x) + \frac{\sigma_t^2}{2}\nabla \log p_t(x)$ and diffusion σ_t. By Theorem 4.14 (sufficiency), it suffices to verify that p_t satisfies the Fokker–Planck equation with this drift:

$$\partial_t p_t(x) = -\mathrm{div}(p_t\, \tilde{u}_t)(x) + \frac{\sigma_t^2}{2}\,\Delta p_t(x). \tag{4.36}$$

Key identity. Before starting, we record one identity that connects the score to the Laplacian. Wherever $p_t(x) > 0$:

$$\nabla p_t = p_t\, \nabla \log p_t \qquad \Longleftrightarrow \qquad \mathrm{div}(p_t\, \nabla \log p_t) = \Delta p_t. \tag{4.37}$$

The left equality is the chain rule ($\nabla \log p_t = \nabla p_t / p_t$). The right follows by substituting: $\mathrm{div}(p_t\, \nabla \log p_t) = \mathrm{div}(\nabla p_t) = \Delta p_t$.

Proof of the SDE Extension Trick (cont.)

Verification. We start from what we *know* and arrive at (4.36).

By the Marginalization Trick (Theorem 4.10), u_t^{target} generates the marginal path p_t via the *continuity equation* (i.e. Fokker–Planck with $\sigma = 0$):

$$\partial_t p_t(x) = -\operatorname{div}(p_t\, u_t^{\text{target}})(x). \tag{4.38}$$

Now expand the RHS of the target equation (4.36):

$$\begin{aligned}
\text{RHS of (4.36)} &= -\operatorname{div}\Big(p_t \left[u_t^{\text{target}} + \tfrac{\sigma_t^2}{2}\nabla \log p_t\right]\Big)(x) + \frac{\sigma_t^2}{2}\Delta p_t(x) \\
&= -\operatorname{div}(p_t\, u_t^{\text{target}})(x) \;-\; \frac{\sigma_t^2}{2}\underbrace{\operatorname{div}(p_t \nabla \log p_t)(x)}_{=\,\Delta p_t(x) \text{ by (4.37)}} + \frac{\sigma_t^2}{2}\Delta p_t(x) \\
&= -\operatorname{div}(p_t\, u_t^{\text{target}})(x) \\
&= \partial_t p_t(x). \qquad \text{(by (4.38))}
\end{aligned}$$

So p_t satisfies (4.36). Since $p_0 = p_{\text{init}}$, Theorem 4.14 (sufficiency) gives $X_t \sim p_t$ for all $t \in [0,1]$. □

The SDE Extension Trick — Why We Need a Conditional Version

The problem with the marginal version. Theorem 4.15 tells us the correct SDE drift is $u_t^{\text{target}}(x) + \frac{\sigma_t^2}{2} \nabla \log p_t(x)$. But *neither* term is tractable:

- The marginal vector field $u_t^{\text{target}}(x) = \int u_t^{\text{target}}(x \mid z) \, \frac{p_t(x|z)\, p_{\text{data}}(z)}{p_t(x)} \, \mathrm{d}z$ requires integrating over p_{data}
- The marginal score $\nabla \log p_t(x)$ requires evaluating the marginal density $p_t(x) = \int p_t(x \mid z)\, p_{\text{data}}(z) \, \mathrm{d}z$

Both integrals are intractable. For *training*, we need a version involving only *conditional* quantities, which we can evaluate in closed form.

The SDE Extension Trick — Conditional Version

Conditional SDE Extension Trick. For each fixed data point z, the SDE

$$\mathrm{d}X_t = \left[u_t^{\text{target}}(X_t \mid z) + \frac{\sigma_t^2}{2} \nabla \log p_t(X_t \mid z) \right] \mathrm{d}t + \sigma_t \, \mathrm{d}W_t, \qquad X_0 \sim p_{\text{init}},$$

satisfies $X_t \sim p_t(\cdot \mid z)$. The proof is identical: apply Fokker–Planck to $p_t(\cdot \mid z)$ in place of p_t, using the fact that $u_t^{\text{target}}(\cdot \mid z)$ generates $p_t(\cdot \mid z)$ via the continuity equation.

For Gaussian paths $p_t(\cdot \mid z) = \mathcal{N}(\alpha_t z, \beta_t^2 I_d)$, both quantities are known in closed form (equations (4.23) and (4.29)):

$$u_t^{\text{target}}(x \mid z) = \left(\dot{\alpha}_t - \frac{\dot{\beta}_t}{\beta_t} \alpha_t \right) z + \frac{\dot{\beta}_t}{\beta_t} x, \qquad \nabla \log p_t(x \mid z) = -\frac{x - \alpha_t z}{\beta_t^2}.$$

These are the training targets.

Langevin Dynamics as a Special Case

Setup. Consider the SDE Extension Trick (Theorem 4.15) in a degenerate case: the probability path is *static*, i.e. $p_t = p$ for all t.

- Since the distribution is not changing, no transport is needed: $u_t^{\text{target}} = 0$
- (Check: the continuity equation $\partial_t p = -\text{div}(p\, u^{\text{target}})$ is satisfied trivially)

The SDE (4.35) reduces to:

$$\mathrm{d}X_t = \frac{\sigma_t^2}{2} \nabla \log p(X_t)\, \mathrm{d}t + \sigma_t\, \mathrm{d}W_t. \tag{4.39}$$

This is the Langevin SDE. By the extension trick: if $X_0 \sim p$, then $X_t \sim p$ for all $t \geq 0$. In other words, p is a stationary distribution of the SDE.

Convergence to equilibrium. Even if $X_0 \not\sim p$, under mild conditions $X_t \to p$ as $t \to \infty$. So Langevin dynamics provides a way to *sample* from any distribution p, given only access to its score $\nabla \log p$.

Langevin Dynamics vs. Generative Modelling

Sampling from a known density. Langevin dynamics (4.39) solves the following problem: given a density p that you can *evaluate* (up to a normalising constant), produce samples $X \sim p$ by running the SDE using $\nabla \log p$.

This is a classical task in computational statistics (MCMC). Crucially, it assumes *analytic access* to p — you can compute $\nabla \log p(x)$ at any query point x.

Generative modelling is a different problem. We do *not* have access to the density p_{data}. We only have a finite dataset $z_1, \ldots, z_N \sim p_{\text{data}}$. We cannot evaluate $\nabla \log p_{\text{data}}(x)$ — we do not even know p_{data} as a formula.

The bridge. This is precisely why we built the machinery of conditional probability paths and the Marginalization Trick: they reduce the problem to *conditional* quantities $u_t^{\text{target}}(x \mid z)$ and $\nabla \log p_t(x \mid z)$ that are known in closed form and can be evaluated given a single data point z from the dataset. A neural network is then trained to approximate them.

From Langevin Dynamics to Diffusion Models

Langevin as a generative model? Early score-based models[17] attempted to bridge the gap: learn $\nabla \log p_{\text{data}}$ with a neural network (trained via *score matching*), then run Langevin dynamics to generate samples.

The limitation. Langevin targets a *static* distribution. Convergence to equilibrium can be very slow when p_{data} is concentrated on a low-dimensional manifold, as is typical for images.

The solution. Instead of waiting for a single static SDE to converge, *transport* the distribution along a time-varying path p_t from p_{init} to p_{data} in finite time $t \in [0, 1]$. This is exactly what the SDE Extension Trick does with $u_t^{\text{target}} \neq 0$:

$$\mathrm{d}X_t = \Big[\underbrace{u_t^{\text{target}}(X_t)}_{\text{transport}} + \underbrace{\frac{\sigma_t^2}{2} \nabla \log p_t(X_t)}_{\text{score correction}} \Big] \mathrm{d}t + \sigma_t \, \mathrm{d}W_t.$$

Langevin dynamics is the special case $u_t^{\text{target}} = 0$. Diffusion models[18] use both terms.

[17] Song, Y. & Ermon, S. (2019). Generative modeling by estimating gradients of the data distribution. *NeurIPS 2019*.

[18] Song, Y., Sohl-Dickstein, J., Kingma, D.P., Kumar, A., Ermon, S. & Poole, B. (2021). Score-based generative modeling through stochastic differential equations. *ICLR 2021*.

Summary: Constructing the Training Target (1/2)

Step 1: Conditional probability path. Choose a family $p_t(x \mid z)$ that interpolates between noise and data, one data point at a time:

$$p_0(\cdot \mid z) = p_{\text{init}} = \mathcal{N}(0, I_d) \qquad \xrightarrow{t \in [0,1]} \qquad p_1(\cdot \mid z) = \delta_z.$$

For Gaussian paths: $p_t(\cdot \mid z) = \mathcal{N}(\alpha_t z, \beta_t^2 I_d)$, with schedulers $\alpha_0 = \beta_1 = 0$, $\alpha_1 = \beta_0 = 1$.

Step 2: Conditional vector field. Find a vector field $u_t^{\text{target}}(x \mid z)$ whose ODE transports $p_0(\cdot \mid z)$ to $p_1(\cdot \mid z)$. For Gaussian paths (equation (4.23)):

$$u_t^{\text{target}}(x \mid z) = \left(\dot{\alpha}_t - \frac{\dot{\beta}_t}{\beta_t} \alpha_t \right) z + \frac{\dot{\beta}_t}{\beta_t} x.$$

Step 3: Marginal vector field. The Marginalization Trick (Theorem 4.10) shows that the mixture

$$u_t^{\text{target}}(x) = \mathbb{E}_{z \sim p_{\text{data}}} \left[u_t^{\text{target}}(x \mid z) \, \frac{p_t(x \mid z)}{p_t(x)} \right]$$

transports p_{init} to p_{data} via the ODE $\mathrm{d}X_t = u_t^{\text{target}}(X_t)\,\mathrm{d}t$. This is a *flow model*.

Summary: Constructing the Training Target (2/2)

Step 4: SDE Extension Trick. For any $\sigma_t \geq 0$, adding noise while correcting the drift with the score preserves the same marginal path (Theorem 4.15):

$$\mathrm{d}X_t = \Big[\underbrace{u_t^{\text{target}}(X_t)}_{\text{flow model drift}} + \underbrace{\frac{\sigma_t^2}{2} \nabla \log p_t(X_t)}_{\text{score correction}} \Big] \mathrm{d}t + \sigma_t \, \mathrm{d}W_t.$$

Setting $\sigma_t = 0$ recovers the flow model. Setting $\sigma_t > 0$ gives a *diffusion model*.

Step 5: Conditional version for training. The marginal quantities $u_t^{\text{target}}(x)$ and $\nabla \log p_t(x)$ are intractable. But their conditional counterparts are known in closed form for Gaussian paths (equations (4.23) and (4.29)):

$$u_t^{\text{target}}(x \mid z) = \left(\dot{\alpha}_t - \frac{\dot{\beta}_t}{\beta_t} \alpha_t \right) z + \frac{\dot{\beta}_t}{\beta_t} x, \qquad \nabla \log p_t(x \mid z) = -\frac{x - \alpha_t z}{\beta_t^2}.$$

These are the targets that a neural network $u_t^\theta(x)$ is trained to approximate.

End of Week 4